

Semantic Spotter Project - Fashion Search AI

(*LangChain-based Generative Semantic Search System*)

Prepared by: **Huy Vo Duc**

Course 6 - GenAI Elective – Module 15

upGrad and IIITB Machine Learning & AI Program – November 2024

Contents

1	Background and Motivation	2
2	Objectives	3
3	Why LangChain?	3
4	Dataset Summary	4
5	Overall System Design.....	5
6	Conceptual System Layers	6
6.1	Data and Embedding Layer	6
6.2	Semantic Retrieval and Ranking Layer	6
6.3	Generative Recommendation Layer.....	6
7	Implementation Details	7
8	Challenges Faced	8
9	Lessons Learned.....	8
10	Conclusion	9

1 Background and Motivation

In previous assignments, semantic search systems were implemented using a custom, manual pipeline, where individual components such as embedding generation, vector similarity search, reranking logic, and response formatting were coded explicitly. While this approach helped build a strong conceptual understanding of retrieval augmented generation (RAG) systems, it introduced several challenges related to scalability, maintainability, and extensibility.

As modern LLM-based applications grow in complexity, managing these components manually can become error-prone and difficult to evolve. This motivates the adoption of higher-level frameworks such as LangChain, which provide abstractions for building robust, modular, and LLM-agnostic generative systems.

The objective of this project is to revisit the semantic search problem from a more production-oriented perspective, leveraging LangChain to simplify orchestration, reduce boilerplate code, and improve system clarity, while preserving the core idea of semantic retrieval and explainable recommendations.

2 Objectives

The goal of the **Fashion Search AI** project is to build an end-to-end generative recommendation system that allows users to search for a fashion catalog using natural language queries and receive:

- Relevant product recommendations based on semantic similarity
- Clear, human-readable explanations describing why each product matches the query

The system is designed to reflect a realistic e-commerce search scenario, where users may search using incomplete, vague, or intent-driven phrases (e.g., “*red women dress for party*”), rather than exact product names.

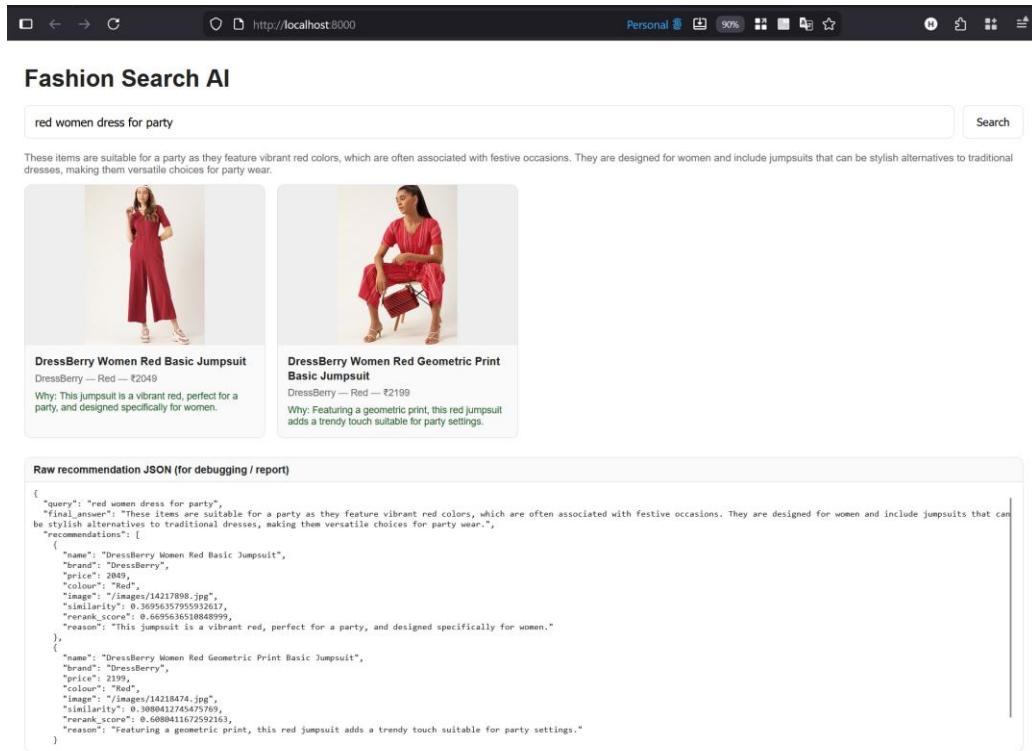


Figure 1: example “red women dress for party”

3 Why LangChain?

LangChain plays a central role in this project and is not used merely as a wrapper library, but as the core orchestration framework.

Key reasons for choosing LangChain:

- **Abstraction of common LLM workflows**
LangChain abstracts repetitive tasks such as embedding models, vector stores, and LLM invocation into reusable building blocks.
- **Seamless vector database integration**
LangChain provides native integration with Chroma, enabling persistent vector stores to be reused across application restarts without manual client management.
- **LLM-agnostic design**
The system can switch between an online LLM (e.g., OpenAI) and an offline rule-based or local generation strategy without changing the retrieval pipeline.
- **Cleaner and more modular codebase**
Compared to the previous assignment, LangChain significantly reduces glue code and improves the readability of the overall system.

Comparison with previous manual implementation:

Aspect	Previous Assignment	This Project (LangChain)
Embedding handling	Manually implemented	Framework-managed embeddings
Vector search	Direct database API calls	LangChain vector store abstraction
LLM integration	Custom prompt + API calls	Prompt templates + LLM chains
Pipeline extensibility	Tightly coupled	Modular and configurable
Production readiness	Limited	Significantly improved

This demonstrates a clear evolution from **concept-focused implementation to a framework-driven, production-ready design**.

4 Dataset Summary

- Source: **Myntra Fashion Dataset** (Kaggle)
- Type: Product metadata and descriptions
- #Rows: ~14,000 usable records
- Product images including
- Price currency: INR (Indian Rupees)

5 Overall System Design

The Fashion Search AI system follows a modular, end-to-end design built using **FastAPI**, **LangChain**, and **Chroma**.

At a high level, the system consists of:

- A backend service exposing a recommendation API
- A persistent semantic vector store for fast retrieval
- A generative component to produce human-readable explanations
- A lightweight web-based user interface for interaction
- For development and tracing so, there is technical panel for displaying technical information

LangChain acts as the **central coordination layer**, connecting embedding models, the vector store, and the LLM-based generation logic into a unified pipeline.

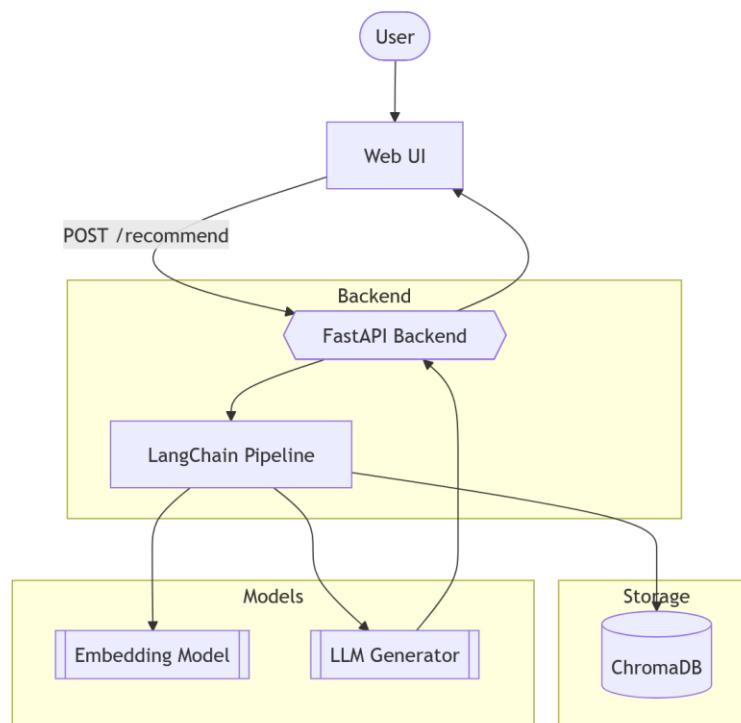


Figure 2: Overall System Architecture of Fashion Search AI

6 Conceptual System Layers

Although the system is presented to users as a single intelligent search experience, it can be conceptually divided into the following layers:

6.1 Data and Embedding Layer

Fashion product descriptions are converted into dense vector embeddings using a Sentence Transformer model (all-MiniLM-L6-v2). These embeddings capture semantic meaning beyond simple keywords and are stored persistently in a Chroma vector database.

6.2 Semantic Retrieval and Ranking Layer

When a user submits a query, LangChain retrieves the most semantically similar products from the vector store. A lightweight hybrid reranking strategy is then applied to improve relevance based on explicit attributes such as color, product type, and gender.

6.3 Generative Recommendation Layer

The top-ranked products are passed to a LangChain-driven generative component. This module produces a structured JSON response that includes:

- a concise summary of the results
- a short explanation for each recommendation explaining its relevance to the user query

This layer improves transparency and user trust by clearly communicating why certain products are recommended.

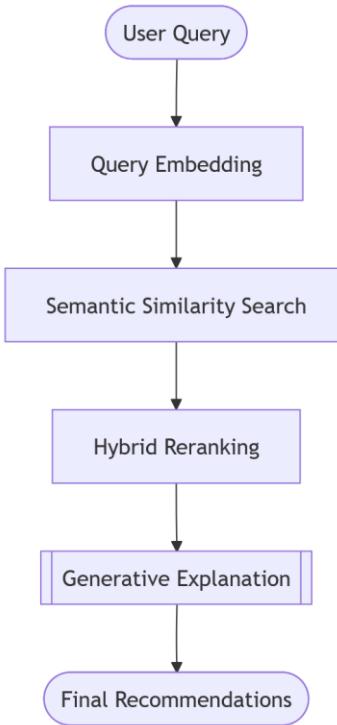


Figure 3: Semantic Spotter Search and Recommendation Workflow

7 Implementation Details

The backend is implemented using **FastAPI** and organized in a clean, modular structure.

- Key implementation aspects include:
- Precomputed embeddings stored in Chroma for fast startup and reuse
- LangChain vector store abstraction for retrieval
- Prompt templates designed to enforce structured JSON output
- A fallback generation mechanism for offline or API-key-free execution
- Environment-based configuration using .env files to avoid hardcoded values

The UI communicates with the backend through a single **/recommend** endpoint, making the system intuitive and aligned with real-world product search experiences.

```
FashionAI/
├── app/
│   ├── main.py
│   ├── routers/
│   │   └── recommend.py
│   ├── services/
│   │   ├── langchain_pipeline.py
│   │   ├── reranker.py
│   │   └── generation.py
│   └── utils/
│       ├── settings.py
│       └── logger.py
├── scripts/
│   └── build_db.py
├── vector_store/
│   └── chroma/
└── ui/
    └── index.html
└── prompts/
    └── prompt_generation.txt
└── README.md
└── requirements.txt
```

Figure 4: Product Structure

8 Challenges Faced

Several challenges were encountered during development:

- Dataset preprocessing required manual validation due to inconsistent records
- Library deprecations in LangChain and Chroma required dependency updates
- Designing prompts that reliably produce valid JSON outputs
- Managing development issues on Windows such as port conflicts and process handling
- Balancing semantic similarity with explicit constraints like color matching

Addressing these challenges led to a deeper understanding of both data-centric and system-level considerations in LLM-based applications.

9 Lessons Learned

This project provided several key takeaways:

- Frameworks like LangChain significantly improve development speed and code clarity
- Semantic search provides a superior user experience compared to keyword-based search

- Explainability is a crucial component of recommendation systems
- Modular system design enables easier experimentation and future extension
- Practical issues such as data quality and dependency management are critical in real-world AI projects

10 Conclusion

Fashion Search AI demonstrates how a modern generative search system can be built using LangChain to combine semantic retrieval with explainable recommendations. By transitioning from a manual pipeline in earlier assignments to a framework-driven approach, the project showcases not only technical competence but also architectural maturity.

The resulting system is robust, extensible, and well-suited for real-world fashion search applications, serving as a strong foundation for further enhancements such as personalization, filtering, and large-scale deployment.